

SMT Solvers: Introduction & Applications

Cambridge 2007

Leonardo de Moura

leonardo@microsoft.com

Microsoft Research

Introduction

- ▶ Industry tools rely on powerful verification engines.
 - ▶ Boolean satisfiability (SAT) solvers.
 - ▶ Binary decision diagrams (BDDs).
- ▶ *Satisfiability Modulo Theories (SMT)*
 - ▶ The next generation of verification engines.
 - ▶ *SAT solvers + Theories*
 - ▶ Arithmetic
 - ▶ Arrays
 - ▶ Uninterpreted Functions
 - ▶ Some problems are more naturally expressed in SMT.
 - ▶ More automation.

Examples

- ▶ $x + 2 = y \Rightarrow f(\text{read}(\text{write}(a, x, 3), y - 2)) = f(y - x + 1)$
- ▶ $f(f(x) - f(y)) \neq f(z), x + z \leq y \leq x \Rightarrow z < 0$

Applications

- ▶ Extended Static Checking.
 - ▶ *Microsoft Spec# and ESP*, ESC/Java
- ▶ Predicate Abstraction.
 - ▶ *Microsoft SLAM/SDV (device driver verification)*.
- ▶ Bounded Model Checking (BMC) & k -induction.
- ▶ Interactive theorem provers (e.g., PVS).
- ▶ Test-case generation.
 - ▶ *Microsoft MUTT*.
- ▶ Symbolic Simulation.
- ▶ Planning & Scheduling.
- ▶ Equivalence checking.

SMT-Solvers & SMT-Lib & SMT-Comp

- ▶ SMT-Solves:

ArgoLib, Ario, Barcellogic, CVC, CVC Lite, CVC3, ExtSAT, Fx7, Harvey, HTP, *ICS (SRI)*, Jat, MathSAT, Sateen, Simplify, Spear, STeP, STP, SVC, TSAT, UCLID, *Yices (SRI)*, Zap (Microsoft), *Z3 (Microsoft)*

- ▶ SMT-Lib: library of benchmarks

`http://www.smtlib.org/`

- ▶ SMT-Comp: annual SMT-Solver competition

`http://www.smtcomp.org/`

Roadmap

- ▶ Background
- ▶ Theories
- ▶ Combination of Theories
- ▶ Quantifiers
- ▶ Applications

Language

▶ A *signature* Σ is a finite set of: function symbols $\Sigma_F = \{f, g, \dots\}$, predicate symbols $\Sigma_P = \{p, q, \dots\}$, and an *arity* function $\Sigma \mapsto N$.

▶ Function symbols with arity 0 are called *constants*.

▶ A countable set \mathcal{V} of *variables* $\{x, y, \dots\}$ disjoint of Σ .

▶ *Terms*:

$$t := f(t_1, \dots, t_n) \mid x$$

▶ *Formulas*:

$$\phi := p(t_1, \dots, t_n) \mid \phi_1 \vee \phi_2 \mid \phi_1 \wedge \phi_2 \mid \neg \phi_1 \mid \exists x : \phi_1 \mid \forall x : \phi_1$$

▶ *Free* (occurrences) of *variables* in a formula are those not bound by a quantifier.

▶ A *sentence* is a first-order formula with no free variables.

Theories

- ▶ A *(first-order) theory* \mathcal{T} (over a signature Σ) is a set of (deductively closed) sentences (over Σ and \mathcal{V}).
- ▶ Let $DC(\Gamma)$ be the deductive closure of a set of sentences Γ .
 - ▶ For every theory \mathcal{T} , $DC(\mathcal{T}) = \mathcal{T}$.
- ▶ A theory \mathcal{T} is *consistent* if *false* $\notin \mathcal{T}$.
- ▶ We can view a (first-order) theory \mathcal{T} as the class of all *models* of \mathcal{T} (due to completeness of first-order logic).

Models (Semantics)

- ▶ A model M is defined as:
 - ▶ Domain S : set of elements.
 - ▶ Interpretation $f^M : S^n \mapsto S$ for each $f \in \Sigma_F$ with $\text{arity}(f) = n$.
 - ▶ Interpretation $p^M \subseteq S^n$ for each $p \in \Sigma_P$ with $\text{arity}(p) = n$.
 - ▶ Assignment $x^M \in S$ for every variable $x \in \mathcal{V}$.
- ▶ A formula ϕ is true in a model M if it evaluates to true under the given interpretations over the domain S .
- ▶ M is a *model for the theory* \mathcal{T} if all sentences of \mathcal{T} are true in M .

Satisfiability and Validity

- ▶ A formula $\phi(\vec{x})$ is *satisfiable* in a theory \mathcal{T} if there is a model of $DC(\mathcal{T} \cup \exists \vec{x}.\phi(\vec{x}))$. That is, there is a model M for \mathcal{T} in which $\phi(\vec{x})$ evaluates to true, denoted by,

$$M \models_{\mathcal{T}} \phi(\vec{x})$$

- ▶ This is also called *\mathcal{T} -satisfiability*.
- ▶ A formula $\phi(\vec{x})$ is *valid* in a theory \mathcal{T} if $\forall \vec{x}.\phi(\vec{x}) \in \mathcal{T}$. That is $\phi(\vec{x})$ evaluates to true in every model M of \mathcal{T} .
- ▶ *\mathcal{T} -validity* is denoted by $\models_{\mathcal{T}} \phi(\vec{x})$.
- ▶ The *quantifier free \mathcal{T} -satisfiability problem* restricts ϕ to be *quantifier free*.

Roadmap

- ▶ Background
- ▶ Theories
- ▶ Combination of Theories
- ▶ Quantifiers
- ▶ Applications

Pure Theory of Equality (EUF)

- ▶ The theory $\mathcal{T}_{\mathcal{E}}$ of equality is the theory $DC(\emptyset)$.
- ▶ The exact set of sentences of $\mathcal{T}_{\mathcal{E}}$ depends on the *signature* in question.
- ▶ The theory does not restrict the possible values of the symbols in its signature in any way. For this reason, it is sometimes called the theory of *equality and uninterpreted functions*.
- ▶ The satisfiability problem for $\mathcal{T}_{\mathcal{E}}$ is the satisfiability problem for first-order logic, which is undecidable.
- ▶ The satisfiability problem for conjunction of literals in $\mathcal{T}_{\mathcal{E}}$ is decidable in polynomial time using *congruence closure*.

Linear Integer Arithmetic

- ▶ $\Sigma_P = \{\leq\}$, $\Sigma_F = \{0, 1, +, -\}$.
- ▶ Let M_{LIA} be the standard model of integers.
- ▶ Then \mathcal{T}_{LIA} is defined to be the set of all Σ sentences true in the model M_{LIA} .
- ▶ As showed by Presburger, the general satisfiability problem for \mathcal{T}_{LIA} is decidable, but its complexity is triply-exponential.
- ▶ The quantifier free satisfiability problem is NP-complete.
- ▶ Remark: non-linear integer arithmetic is undecidable even for the quantifier free case.

Linear Real Arithmetic

- ▶ The general satisfiability problem for $\mathcal{T}_{\mathcal{LRA}}$ is decidable, but its complexity is doubly-exponential.
- ▶ The quantifier free satisfiability problem is solvable in polynomial time, though exponential methods (Simplex) tend to perform best in practice.

Difference Logic

- ▶ *Difference logic* is a fragment of linear arithmetic.
- ▶ Atoms have the form: $x - y \leq c$.
- ▶ Most linear arithmetic atoms found in hardware and software verification are in this fragment.
- ▶ The quantifier free satisfiability problem is solvable in $O(nm)$.

Theory of Arrays

▶ $\Sigma_P = \emptyset, \Sigma_F = \{read, write\}.$

▶ Non-extensional arrays

▶ Let $\Lambda_{\mathcal{A}}$ be the following axioms:

$$\forall a, i, v. read(write(a, i, v), i) = v$$

$$\forall a, i, j, v. i \neq j \Rightarrow read(write(a, i, v), j) = read(a, j)$$

▶ $\mathcal{T}_{\mathcal{A}} = DC(\Lambda_{\mathcal{A}})$

▶ For extensional arrays, we need the following extra axiom:

$$\forall a, b. (\forall i. read(a, i) = read(b, i)) \Rightarrow a = b$$

▶ The satisfiability problem for $\mathcal{T}_{\mathcal{A}}$ is undecidable, the quantifier free case is NP-complete.

Theory of Bit-vectors

- ▶ Bit-vectors (also called “words”) are a crucial abstraction for reasoning about hardware and software structures.
- ▶ Operations:
 - ▶ Concatenation.
 - ▶ Extraction.
 - ▶ Bit-wise operations.
 - ▶ Modular arithmetic.
- ▶ Deciding equality of arbitrary combinations of these operations is NP-hard.

Transitive closure

- ▶ The *transitive closure* R^* of a binary relation R on a set X is the smallest transitive relation on X that contains R .
- ▶ Useful for reasoning about programs that manipulate heap objects (e.g., lists and trees).
- ▶ For example, if X is a set of heap objects, and $R(x, y)$ is the relation 'x points to y', then $R^*(x, y)$ is the relation 'x reaches y'.

Other theories

- ▶ Partial orders
- ▶ Tuples & Records
- ▶ Inductive datatypes
- ▶ ...

Roadmap

- ▶ Background
- ▶ Theories
- ▶ Combination of Theories
- ▶ Quantifiers
- ▶ Applications

Combination of Theories

- ▶ In practice, we need a combination of theories.

- ▶ Examples:

- ▶ $x + 2 = y \Rightarrow f(\text{read}(\text{write}(a, x, 3), y - 2)) = f(y - x + 1)$

- ▶ $f(f(x) - f(y)) \neq f(z), x + z \leq y \leq x \Rightarrow z < 0$

- ▶ Given

$$\Sigma = \Sigma_1 \cup \Sigma_2$$

$$\mathcal{T}_1, \mathcal{T}_2 : \text{theories over } \Sigma_1, \Sigma_2$$

$$\mathcal{T} = DC(\mathcal{T}_1 \cup \mathcal{T}_2)$$

- ▶ Is \mathcal{T} consistent?
- ▶ Given satisfiability procedures for conjunction of literals of \mathcal{T}_1 and \mathcal{T}_2 , how to decide the satisfiability of \mathcal{T} ?

Preamble

- ▶ Disjoint signatures: $\Sigma_1 \cap \Sigma_2 = \emptyset$.
- ▶ Stably-Infinite Theories.
- ▶ Convex Theories.

Stably-Infinite Theories

- ▶ A theory is *stably infinite* if every satisfiable QFF is satisfiable in an infinite model.
- ▶ Example. Theories with only finite models are not stably infinite.
$$\mathcal{T}_2 = DC(\forall x, y, z. (x = y) \vee (x = z) \vee (y = z)).$$
- ▶ Is this a problem in practice? (We want to support the “finite types” found in our programming languages)

Stably-Infinite Theories

- ▶ A theory is *stably infinite* if every satisfiable QFF is satisfiable in an infinite model.
- ▶ Example. Theories with only finite models are not stably infinite.
 $\mathcal{T}_2 = DC(\forall x, y, z. (x = y) \vee (x = z) \vee (y = z)).$
- ▶ Is this a problem in practice? (We want to support the “finite types” found in our programming languages)
- ▶ Answer: *No*. \mathcal{T}_2 is not useful in practice. Add a predicate $in_2(x)$ (intuition: x is an element of the “finite type”).

$$\mathcal{T}_2' = DC(\forall x, y, z. in_2(x) \wedge in_2(y) \wedge in_2(z) \Rightarrow (x = y) \vee (x = z) \vee (y = z))$$

- ▶ \mathcal{T}_2' is *stably infinite*.

Stably-Infinite Theories (cont.)

- ▶ *The union of two consistent, disjoint, stably infinite theories is consistent.*

Convexity

- ▶ A theory \mathcal{T} is *convex* iff
 - for all finite sets Γ of literals and
 - for all non-empty disjunctions $\bigvee_{i \in I} x_i = y_i$ of variables,
 $\Gamma \models_{\mathcal{T}} \bigvee_{i \in I} x_i = y_i$ iff $\Gamma \models_{\mathcal{T}} x_i = y_i$ for some $i \in I$.
- ▶ Every convex theory \mathcal{T} with non trivial models (i.e., $\models_{\mathcal{T}} \exists x, y. x \neq y$) is stably infinite.
- ▶ All *Horn* theories are convex – this includes all (conditional) equational theories.
- ▶ *Linear rational arithmetic is convex.*

Convexity (cont.)

▶ *Many theories are not convex:*

▶ Linear integer arithmetic.

$$y = 1, z = 2, 1 \leq x \leq 2 \models x = y \vee x = z$$

▶ Nonlinear arithmetic.

$$x^2 = 1, y = 1, z = -1 \models x = y \vee x = z$$

▶ Theory of Bit-vectors.

▶ Theory of Arrays.

$$v_1 = \text{read}(\text{write}(a, i, v_2), j), v_3 = \text{read}(a, j) \models \\ v_1 = v_2 \vee v_1 = v_3$$

Convexity: Example

- ▶ Let $\mathcal{T} = \mathcal{T}_1 \cup \mathcal{T}_2$, where \mathcal{T}_1 is EUF ($O(n \log(n))$) and \mathcal{T}_2 is IDL ($O(nm)$).
- ▶ \mathcal{T}_2 is not convex.
- ▶ Satisfiability is NP-Complete for $\mathcal{T} = \mathcal{T}_1 \cup \mathcal{T}_2$.
 - ▶ Reduce 3CNF satisfiability to \mathcal{T} -satisfiability.
 - ▶ For each boolean variable p_i add the atomic formulas:
 $0 \leq x_i, x_i \leq 1$.
 - ▶ For a clause $p_1 \vee \neg p_2 \vee p_3$ add the atomic formula:
 $f(x_1, x_2, x_3) \neq f(0, 1, 0)$

Nelson-Oppen Combination

- Let \mathcal{T}_1 and \mathcal{T}_2 be consistent, stably infinite theories over disjoint (countable) signatures. Assume satisfiability of conjunction of literals can be decided in $O(T_1(n))$ and $O(T_2(n))$ time respectively.

Then,

1. The combined theory \mathcal{T} is consistent and stably infinite.
2. Satisfiability of quantifier free conjunction of literals in \mathcal{T} can be decided in $O(2^{n^2} \times (T_1(n) + T_2(n)))$.
3. If \mathcal{T}_1 and \mathcal{T}_2 are convex, then so is \mathcal{T} and satisfiability in \mathcal{T} is in $O(n^3 \times (T_1(n) + T_2(n)))$.

Reduction Functions

- ▶ A *reduction function* reduces the satisfiability problem for a theory \mathcal{T}_1 to the satisfiability problem of a simpler theory \mathcal{T}_2 .
- ▶ Reduction functions simplify the implementation.
- ▶ Potential disadvantages:
 - ▶ “Information loss”.
 - ▶ Eager addition of irrelevant information.
- ▶ Theory of commutative functions.
 - ▶ Deductive closure of: $\forall x, y. f(x, y) = f(y, x)$
 - ▶ Reduction to $\mathcal{T}_{\mathcal{E}}$.
 - ▶ For every $f(a, b)$ in ϕ , add the equality $f(a, b) = f(b, a)$.

Reduction Functions: Ackermann's reduction

- ▶ *Ackermann's reduction* is used to remove uninterpreted functions.
 - ▶ For each application $f(\vec{a})$ in ϕ create a fresh variable $f_{\vec{a}}$.
 - ▶ For each pair of applications $f(\vec{a}), f(\vec{c})$ in ϕ add the clause $\vec{a} \neq \vec{c} \vee f_{\vec{a}} = f_{\vec{c}}$.
 - ▶ Replace $f(\vec{a})$ with $f_{\vec{a}}$ in ϕ .
- ▶ It is used in some SMT solvers to reduce $\mathcal{T}_{\mathcal{LA}} \cup \mathcal{T}_{\mathcal{E}}$ to $\mathcal{T}_{\mathcal{LA}}$.
- ▶ *Main problem: quadratic number of new clauses.*
- ▶ It is also problematic to use this approach in the context of *several theories* and when combining SMT solvers with *quantifier instantiation*.

Breakthrough in SAT solving

- ▶ Breakthrough in SAT solving influenced the way SMT solvers are implemented.
- ▶ Modern SAT solvers are based on the DPLL algorithm.
- ▶ Modern implementations add several sophisticated *search techniques*.
 - ▶ Backjumping
 - ▶ Learning
 - ▶ Restarts
 - ▶ Watched literals

The Eager Approach

- ▶ Translate formula into equisatisfiable propositional formula and use off-the-shelf SAT solver.
- ▶ Why “eager”?
Search uses *all* theory information from the beginning.
- ▶ Can use best available SAT solver.
- ▶ Sophisticated encodings are need for each theory.
- ▶ Sometimes translation and/or solving too slow.

Lazy approach: SAT solvers + Theories

- ▶ This approach was independently developed by several groups: CVC (Stanford), ICS (SRI), MathSAT (Univ. Trento, Italy), and Verifun (HP).
- ▶ It was motivated also by the breakthroughs in SAT solving.
- ▶ SAT solver “manages” the boolean structure, and assigns truth values to the atoms in a formula.
- ▶ Efficient theory solvers is used to validate the (partial) assignment produced by the SAT solver.
- ▶ When theory solver detects unsatisfiability \rightarrow a new clause (*lemma*) is created.

Roadmap

- ▶ Background
- ▶ Theories
- ▶ Combination of Theories
- ▶ Quantifiers
- ▶ Applications

Quantifiers

- ▶ Since first-order logic is undecidable, satisfiability is not solvable for arbitrary quantified formulas.
- ▶ Some theories, e.g., datatypes, linear arithmetic over integers, arithmetic over reals, support quantifier elimination.
- ▶ Existential quantifiers can be skolemized, but the problem of instantiating universal quantifiers for detecting unsatisfiability remains.

Heuristic Quantifier Instantiation

- ▶ Semantically, $\forall x_1, \dots, x_n. F$ is equivalent to the infinite conjunction $\bigwedge_{\beta} \beta(F)$.
- ▶ Solvers use heuristics to select from this infinite conjunction those instances that are “relevant”.
- ▶ The key idea is to treat an instance $\beta(F)$ as relevant whenever it contains enough terms that are represented in the solver state.
- ▶ Non ground terms p from F are selected as *patterns*.
- ▶ *E-matching* (matching modulo equalities) is used to find instances of the patterns.
- ▶ Example: $f(a, b)$ matches the pattern $f(g(x), x)$ if $a = g(b)$.

E-matching

- ▶ E-matching is NP-hard.
- ▶ The number of matches can be exponential.
- ▶ It is not refutationally complete.
- ▶ In practice:
 - ▶ Indexing techniques for fast retrieval.
 - ▶ Incremental E-matching.

E-matching: example

- ▶ $\forall x. f(g(x)) = x$
- ▶ Pattern: $f(g(x))$
- ▶ Atoms: $a = g(b), b = c, f(a) \neq c$
- ▶ $\rightarrow \text{instantiate } f(g(b)) = b$

E-matching limitations

- ▶ E-matching needs ground (seed) terms.
 - ▶ It fails to prove simple properties when ground (seed) terms are not available.
 - ▶ Example:

$$(\forall x. f(x) \leq 0) \wedge (\forall x. f(x) > 0)$$

- ▶ Matching loops

$$(\forall x. f(x) = g(f(x))) \wedge (\forall x. g(x) = f(g(x)))$$

- ▶ Inefficiency and/or non-termination.
- ▶ Some solvers have support for detecting matching loops based on instantiation chain length.

Quantifiers: future work

- ▶ When quantifiers are used, solvers such as Yices and Z3 produce “candidate models”.

- ▶ Model checking.

Evaluate quantifiers using the “candidate model”, use violations to create new quantifier instantiations.

- ▶ Superposition calculus + SMT.
- ▶ Support for decidable fragments.

Roadmap

- ▶ Background
- ▶ Theories
- ▶ Combination of Theories
- ▶ Quantifiers
- ▶ Applications

Bounded Model Checking (BMC)

- ▶ To check whether a program with initial state I and next-state relation T violates the invariant Inv in the first k steps, one checks:

$$I(s_0) \wedge T(s_0, s_1) \wedge \dots \wedge T(s_{k-1}, s_k) \wedge (\neg Inv(s_0) \vee \dots \vee \neg Inv(s_k))$$

- ▶ This formula is satisfiable if and only if there exists a path of length at most k from the initial state s_0 which violates the invariant I .
- ▶ Formulas produced in BMC are usually quite big.
- ▶ The SAL bounded model checker from SRI uses SMT solvers.
<http://sal.csl.sri.com>

MUTT: MSIL Unit Testing Tools

- ▶ `http://research.microsoft.com/projects/mutt`
- ▶ *Unit tests are popular*, but it is far from trivial to write them.
- ▶ It is quite laborious to write enough of them to have confidence in the correctness of an implementation.
- ▶ Approach: *symbolic execution*.
- ▶ Symbolic execution builds a path condition over the input symbols.
- ▶ A *path condition* is a mathematical formula that encodes data constraints that result from executing a given code path.

MUTT: MSIL Unit Testing Tools

- ▶ When symbolic execution reaches a if-statement, it will explore two execution paths:
 1. The if-condition is conjoined to the path condition for the then-path.
 2. The negated condition to the path condition of the else-path.
- ▶ SMT solver must be able to produce models.
- ▶ SMT solver is also used to test path *feasibility*.

Spec#: Extended Static Checking

- ▶ <http://research.microsoft.com/specsharp/>
- ▶ Superset of C#
 - ▶ non-null types
 - ▶ pre- and postconditions
 - ▶ object invariants
- ▶ Static program verification
- ▶ Example:

```
public StringBuilder Append(char[] value, int startIndex,
                           int charCount);
requires value == null ==> startIndex == 0 && charCount == 0;
requires 0 <= startIndex;
requires 0 <= charCount;
requires value == null ||
        startIndex + charCount <= value.Length;
```

Spec#: Architecture

- ▶ Verification condition generation:

Spec# compiler: Spec# \rightsquigarrow MSIL (bytecode).

Bytecode translator: MSIL \rightsquigarrow Boogie PL.

V.C. generator: Boogie PL \rightsquigarrow SMT formula.

- ▶ SMT solver is used to prove the verification conditions.
- ▶ Counterexamples are traced back to the source code.
- ▶ *The formulas produced by Spec# are not quantifier free.*
 - ▶ Heuristic quantifier instantiation is used.

SLAM: device driver verification

- ▶ <http://research.microsoft.com/slam/>
- ▶ *SLAM/SDV* is a software model checker.
- ▶ Application domain: *device drivers*.
- ▶ Architecture
 - c2bp** C program \rightsquigarrow boolean program (*predicate abstraction*).
 - bebop** Model checker for boolean programs.
 - newton** Model refinement (*check for path feasibility*)
- ▶ SMT solvers are used to perform predicate abstraction and to check path feasibility.
- ▶ c2bp makes several calls to the SMT solver. The formulas are relatively small.

Interactive Theorem Provers

- ▶ *SMT solvers can be used inside of an interactive theorem provers.*
- ▶ More automation.
- ▶ In PVS, goals can be discharged using Yices.
- ▶ Model generation: improved “quick-check”.

Scheduling

- ▶ Given j jobs and m machines, each job consists of a sequence of tasks t_{i1}, \dots, t_{in} , where each task t_{ik} is pair $\langle M, \delta \rangle$ for machine M and duration δ .
- ▶ Find a schedule with a minimum duration.
- ▶ Incremental SMT solving: binary search in $[0, \sum_i \delta_i]$.
- ▶ Example:

Jobs	Tasks
a	$\langle 1, 2 \rangle, \langle 2, 6 \rangle$
b	$\langle 2, 5 \rangle, \langle 1, 3 \rangle, \langle 2, 3 \rangle$
c	$\langle 2, 4 \rangle$
d	$\langle 1, 5 \rangle, \langle 2, 2 \rangle$

Planning

- ▶ Given c cities, t trucks each located at a specific city, and p packages each with source city and destination city.
- ▶ In each step, packages can be loaded and unloaded, or the trucks can be driven from one city to another.
- ▶ Find a plan with a minimum number of steps for delivering the packages from source to destination.
- ▶ For each step i , we have predicates:
 $location(t, c, i)$, $at(p, c, i)$, $on(p, t, i)$.
- ▶ Constraints assert that a package can be either on one truck or at a city, a package can be loaded or unloaded from a truck to a city only if the truck is at the city, etc.

Conclusion

- ▶ SMT is the next generation of verification engines.
- ▶ More automation: it is push-button technology.
- ▶ SMT solvers are used in different applications.
- ▶ The breakthrough in SAT solving influenced the new generation of SMT solvers.