

Z3: An Efficient SMT Solver

MSR Cambridge 2007

Leonardo de Moura and Nikolaj Bjørner

{leonardo, nbjorner}@microsoft.com.

Microsoft Research

Introduction

- ▶ *Z3 is a new SMT solver developed at Microsoft Research.*
- ▶ It is still under development.
- ▶ Version 0.1 is the first external release.
- ▶ New external releases are coming soon.
- ▶ *It replaced Simplify as the default prover for Spec#/Boogie.*
- ▶ It won 4 divisions in SMT-COMP'07.

Satisfiability Modulo Theories (SMT)

- ▶ SMT is the problem of determining *satisfiability* of formulas *modulo* background *theories*.
- ▶ Examples of background theories:
 - ▶ linear arithmetic: $x + y \leq 10$
 - ▶ bit-vectors: $x \& (x - 1)$
 - ▶ uninterpreted functions: $g(f(x)) = x$
 - ▶ arrays: $read(write(a, i, v), j) \neq v$
- ▶ Example of formula:

$$x + 2 = y, f(read(write(a, x, 3), y - 2)) \neq f(y - x + 1)$$

Applications

- ▶ Extended Static Checking.
 - ▶ *Microsoft Spec# and ESP*, ESC/Java
- ▶ Predicate Abstraction.
 - ▶ *Microsoft SLAM/SDV (device driver verification)*.
- ▶ Bounded Model Checking (BMC) & k -induction.
- ▶ Test-case generation.
 - ▶ *Microsoft MUTT*.
- ▶ Symbolic Simulation.
- ▶ Planning & Scheduling.
- ▶ Equivalence checking.

Supported Features

- ▶ Linear real and integer arithmetic.
- ▶ Fixed-size bit-vectors.
- ▶ Uninterpreted Functions.
- ▶ Extensional Arrays.
- ▶ Quantifiers.
- ▶ Model generation.
- ▶ Coming soon:
 - ▶ Improved support for non linear arithmetic.
 - ▶ Improved SAT solver.
 - ▶ Improved support for quantifiers.
 - ▶ Sets & Transitive Closure.

Architecture

- ▶ A modern DPLL-based SAT solver.
- ▶ A *core theory solver* that handles equalities and uninterpreted functions.
- ▶ *Satellite solvers* (for arithmetic, arrays, etc).
- ▶ An *E-matching abstract machine* (for quantifiers).
- ▶ Very modular: new theories can be added without modifying the core.

Z3: new features

- ▶ Model based theory combination.
- ▶ E-matching abstract machine.
- ▶ “Don’t care” propagation.
- ▶ Create/Delete clauses/atoms/terms during backtracking search.

Theory combination

- ▶ Theory Combination Problem.
 - ▶ $T = T_1 \cup T_2$
 - ▶ Is T consistent?
 - ▶ Given decision procedures for T_1 and T_2 , how to decide T ?
- ▶ Nelson-Oppen Approach:
 - ▶ Disjoint theories.
 - ▶ Each theory propagate (disjunction) of implied equalities.
 - ▶ For convex theories only implied equalities need to be propagated.

Combining theories in practice

- ▶ *Propagate all implied equalities.*
 - ▶ Deterministic Nelson-Oppen.
 - ▶ Complete only for convex theories.
 - ▶ It may be expensive for some theories.
- ▶ *Delayed Theory Combination.*
 - ▶ Nondeterministic Nelson-Oppen.
 - ▶ Create set of interface equalities ($x = y$) between shared variables.
 - ▶ Use SAT solver to guess the partition.
 - ▶ Disadvantage: the number of additional equality literals is quadratic in the number of shared variables.

Model based theory combination

- ▶ Common to these methods is that they are *pessimistic* about which equalities are propagated.

- ▶ *Model-based Theory Combination*

- ▶ *Optimistic approach.*

- ▶ Use a candidate model M_i for one of the theories \mathcal{T}_i and propagate all equalities implied by the candidate model, hedging that other theories will agree.

if $M_i \models \mathcal{T}_i \cup \Gamma_i \cup \{u = v\}$ **then** propagate $u = v$.

- ▶ If not, use backtracking to fix the model.
 - ▶ It is cheaper to enumerate equalities that are implied in a particular model than of all models.

Model based theory combination: Example

$$x = f(y - 1), f(x) \neq f(y), 0 \leq x \leq 1, 0 \leq y \leq 1$$

Purifying

Model based theory combination: Example

$$x = f(z), f(x) \neq f(y), 0 \leq x \leq 1, 0 \leq y \leq 1, z = y - 1$$

Model based theory combination: Example

$\mathcal{T}_{\mathcal{E}}$			$\mathcal{T}_{\mathcal{A}}$	
Literals	Eq. Classes	Model	Literals	Model
$x = f(z)$	$\{x, f(z)\}$	$x^{\mathcal{E}} = *_1$	$0 \leq x \leq 1$	$x^{\mathcal{A}} = 0$
$f(x) \neq f(y)$	$\{y\}$	$y^{\mathcal{E}} = *_2$	$0 \leq y \leq 1$	$y^{\mathcal{A}} = 0$
	$\{z\}$	$z^{\mathcal{E}} = *_3$	$z = y - 1$	$z^{\mathcal{A}} = -1$
	$\{f(x)\}$	$f^{\mathcal{E}} = \{*_1 \mapsto *_4,$		
	$\{f(y)\}$	$*_2 \mapsto *_5,$		
		$\text{else} \mapsto *_1,$		

Assume $x = y$

Model based theory combination: Example

$\mathcal{T}_{\mathcal{E}}$			$\mathcal{T}_{\mathcal{A}}$	
Literals	Eq. Classes	Model	Literals	Model
$x = f(z)$ $f(x) \neq f(y)$ $x = y$	$\{x, y, f(z)\}$ $\{z\}$ $\{f(x), f(y)\}$	$x^{\mathcal{E}} = *_1$ $y^{\mathcal{E}} = *_1$ $z^{\mathcal{E}} = *_2$ $f^{\mathcal{E}} = \{*_1 \mapsto *_3,$ $\quad \text{else} \mapsto *_1\}$	$0 \leq x \leq 1$ $0 \leq y \leq 1$ $z = y - 1$ $x = y$	$x^{\mathcal{A}} = 0$ $y^{\mathcal{A}} = 0$ $z^{\mathcal{A}} = -1$

Unsatisfiable

Model based theory combination: Example

$\mathcal{T}_{\mathcal{E}}$			$\mathcal{T}_{\mathcal{A}}$	
Literals	Eq. Classes	Model	Literals	Model
$x = f(z)$	$\{x, f(z)\}$	$x^{\mathcal{E}} = *_1$	$0 \leq x \leq 1$	$x^{\mathcal{A}} = 0$
$f(x) \neq f(y)$	$\{y\}$	$y^{\mathcal{E}} = *_2$	$0 \leq y \leq 1$	$y^{\mathcal{A}} = 0$
$x \neq y$	$\{z\}$	$z^{\mathcal{E}} = *_3$	$z = y - 1$	$z^{\mathcal{A}} = -1$
	$\{f(x)\}$	$f^{\mathcal{E}} = \{*_1 \mapsto *_4,$	$x \neq y$	
	$\{f(y)\}$	$*_2 \mapsto *_5,$		
		$\text{else} \mapsto *_1\}$		

Backtrack, and assert $x \neq y$.

$\mathcal{T}_{\mathcal{A}}$ model need to be fixed.

Model based theory combination: Example

$\mathcal{T}_{\mathcal{E}}$			$\mathcal{T}_{\mathcal{A}}$	
Literals	Eq. Classes	Model	Literals	Model
$x = f(z)$	$\{x, f(z)\}$	$x^{\mathcal{E}} = *_1$	$0 \leq x \leq 1$	$x^{\mathcal{A}} = 0$
$f(x) \neq f(y)$	$\{y\}$	$y^{\mathcal{E}} = *_2$	$0 \leq y \leq 1$	$y^{\mathcal{A}} = 1$
$x \neq y$	$\{z\}$	$z^{\mathcal{E}} = *_3$	$z = y - 1$	$z^{\mathcal{A}} = 0$
	$\{f(x)\}$	$f^{\mathcal{E}} = \{*_1 \mapsto *_4,$	$x \neq y$	
	$\{f(y)\}$	$*_2 \mapsto *_5,$		
		$\text{else} \mapsto *_1\}$		

Assume $x = z$

Model based theory combination: Example

$\mathcal{T}_{\mathcal{E}}$			$\mathcal{T}_{\mathcal{A}}$	
Literals	Eq. Classes	Model	Literals	Model
$x = f(z)$	$\{x, z, f(x), f(z)\}$	$x^{\mathcal{E}} = *_1$	$0 \leq x \leq 1$	$x^{\mathcal{A}} = 0$
$f(x) \neq f(y)$	$\{y\}$	$y^{\mathcal{E}} = *_2$	$0 \leq y \leq 1$	$y^{\mathcal{A}} = 1$
$x \neq y$	$\{f(y)\}$	$z^{\mathcal{E}} = *_1$	$z = y - 1$	$z^{\mathcal{A}} = 0$
$x = z$		$f^{\mathcal{E}} = \{*_1 \mapsto *_1,$ $\text{else} \mapsto *_3,$	$x \neq y$	
			$x = z$	

Satisfiable

Experimental Results

	#	MathSAT	MathSAT-dtc	Yices	Z3
EufLaArithmetic	52	1851.50 (11)	785.87 (1)	10.45	17.34
Hash	199	520.90	19.39	11.48	6.54
Wisa	256	886.36 (1)	6916.18	4.37	2.78
RandomCoupled	400	517.05	518.15	9516.11 (51)	56.16
RandomDecoupled	500	11989.60 (1)	97.07	19362.40 (51)	41.95
Simple	98	1366.33	7053.98 (29)	2328.63 (53)	1.00
Ackermann	99	228.49 (82)	344.00 (82)	2.99	1.72
Total	1604	17360.23 (95)	15734.64 (112)	31236.43 (155)	127.49

Quantifiers

- ▶ Since first-order logic is undecidable, satisfiability is not solvable for arbitrary quantified formulas.
- ▶ Some theories, e.g., datatypes, linear arithmetic over integers, arithmetic over reals, support quantifier elimination.
- ▶ Existential quantifiers can be skolemized, but the problem of instantiating universal quantifiers for detecting unsatisfiability remains.

Heuristic Quantifier Instantiation

- ▶ Semantically, $\forall x_1, \dots, x_n. F$ is equivalent to the infinite conjunction $\bigwedge_{\beta} \beta(F)$.
- ▶ Solvers use heuristics to select from this infinite conjunction those instances that are “relevant”.
- ▶ The key idea is to treat an instance $\beta(F)$ as relevant whenever it contains enough terms that are represented in the solver state.
- ▶ Non ground terms p from F are selected as *patterns*.
- ▶ *E-matching* (matching modulo equalities) is used to find instances of the patterns.
- ▶ Example: $f(a, b)$ matches the pattern $f(g(x), x)$ if $a = g(b)$.

E-matching

- ▶ E-matching is NP-hard.
- ▶ The number of matches can be exponential.
- ▶ It is not refutationally complete.
- ▶ In practice:
 - ▶ Indexing techniques for fast retrieval.
 - ▶ Incremental E-matching.

E-matching: example

- ▶ $\forall x. f(g(x)) = x$
- ▶ Pattern: $f(g(x))$
- ▶ Atoms: $a = g(b), b = c, f(a) \neq c$
- ▶ $\rightarrow \text{instantiate } f(g(b)) = b$

Quantifiers in Z3

- ▶ Z3 uses a E-matching abstract machine.
 - ▶ Patterns \rightsquigarrow code sequence.
 - ▶ Abstract machine executes the code.
- ▶ *Z3 uses new algorithms that identify matches on E-graphs incrementally and efficiently.*
 - ▶ E-matching code trees.
 - ▶ Inverted path index.
- ▶ Z3 garbage collects clauses, together with their atoms and terms, that were useless in closing branches.

E-matching code trees


- ▶ In practice, there are several similar patterns.
- ▶ *Idea: combine several code sequences in a code tree.*
- ▶ Factor out redundant work.
- ▶ Match several patterns simultaneously.
- ▶ Saturation based theorem provers use a different kind of code tree to implement:
 - ▶ Forward subsumption.
 - ▶ Forward demodulation.

Incremental E-matching

- ▶ Z3 uses a backtracking search.
- ▶ *New terms are created during the search.*
 - ▶ A code tree for each function symbol f .
Patterns that start with a f -application.
 - ▶ Execute code-tree for each new term.
- ▶ *New equalities are asserted during the search.*
 - ▶ New equalities \rightsquigarrow new E-matching instances.
 - ▶ Example:
 $f(a, b)$ matches $f(g(x), x)$ after
 $a = g(b)$ is asserted.

Inverted path index

- ▶ It is used to find which patterns may have new instances after an equality is asserted.
- ▶ Inverted path index for *pc-pair* (f, g) and patterns $f(f(g(x), a), x)$, $h(c, f(g(y), x))$, $f(f(g(x), b), y)$, $f(f(a, g(x)), g(y))$.



E-matching limitations

- ▶ E-matching needs ground (seed) terms.
 - ▶ It fails to prove simple properties when ground (seed) terms are not available.
 - ▶ Example:


$$(\forall x. f(x) \leq 0) \wedge (\forall x. f(x) > 0)$$

- ▶ Matching loops


$$(\forall x. f(x) = g(f(x))) \wedge (\forall x. g(x) = f(g(x)))$$

- ▶ Inefficiency and/or non-termination.
- ▶ Some solvers have support for detecting matching loops based on instantiation chain length.

Experimental Results: SMT-LIB




() Z3 vs. CVC3 1.0




() Z3 vs. Yices 1.0

Experimental Results: ESC/Java




() Z3 vs. Simplify




() Z3 vs. Zap 2.0

Experimental Results: Boogie




() Z3 vs. Simplify



() Z3 vs. Zap 2.0

SMT-COMP'07: AUFLIA



Don't cares

- ▶ DPLL(T) based solvers assign a boolean value to potentially all atoms appearing in a goal.
- ▶ In practice, several of these atoms are *don't cares*.
- ▶ *Z3 ignores don't care atoms for expensive inference rules and theories, such as, quantifier instantiation.*

Bit-vector theory

- ▶ Simple approach:
 - ▶ Based on bit-blasting.
 - ▶ Partial evaluation.
 - ▶ “Don’t cares”.
- ▶ *Supports all operations used in software and hardware.*
- ▶ They can be used with other theories and quantifiers.
- ▶ *Z3 uses an efficient and modular pre-processor.*

Getting Z3

- ▶ Internal

- ▶ Stable builds: `http://codebox/z3`

- ▶ Fresh builds: `\\pexbuild3\drops\z3\latest\release`

- ▶ External

- ▶ Official website:

- `http://research.microsoft.com/projects/z3`

- ▶ SMT-COMP'07: `http://www.smtcomp.org`

- ▶ Managed & Unmanaged APIs.

- ▶ SMT-LIB & Simplify input formats.

Current (and future) users

- ▶ Boogie/Spec#
- ▶ SLAM/SDV
- ▶ Pex
- ▶ Vigilante
- ▶ *Your project?*

Future work

- ▶ Improved support for quantifiers:
 - ▶ Model checking.
 - ▶ Superposition calculus + SMT.
 - ▶ Decidable fragments.
- ▶ Machine learning & dynamic tuning.
- ▶ Transitive Closure.
- ▶ Improved bit-vector and array theories.
- ▶ Better performance.
- ▶ More “customers” & applications.

Conclusion

- ▶ Z3 is an efficient and modular SMT solver.
- ▶ It is going to be used in several projects at Microsoft.
- ▶ It can solve 99.7% of the benchmarks in SMT-LIB.
- ▶ New releases of Z3 will be available at:
`http://research.microsoft.com/projects/z3`.