

SMT@Microsoft

Intel 2007

Leonardo de Moura and Nikolaj Bjørner

{leonardo, nbjorner}@microsoft.com.

Microsoft Research

Introduction

- ▶ Industry tools rely on powerful verification engines.
 - ▶ Boolean satisfiability (SAT) solvers.
 - ▶ Binary decision diagrams (BDDs).
- ▶ *Satisfiability Modulo Theories (SMT)*
 - ▶ The next generation of verification engines.
 - ▶ *SAT solvers + Theories*
 - ▶ Arithmetic
 - ▶ Arrays
 - ▶ Uninterpreted Functions
 - ▶ Some problems are more naturally expressed in SMT.
 - ▶ More automation.

SMT: Examples

$$x + 2 = y \Rightarrow f(\text{read}(\text{write}(a, x, 3), y - 2)) = f(y - x + 1)$$

$$f(f(x) - f(y)) \neq f(z), x + z \leq y \leq x \Rightarrow z < 0$$

SMT-Solvers & SMT-Lib & SMT-Comp

- ▶ SMT-Solves:

Ario, Barcellogic, CVC, CVC Lite, CVC3, ExtSAT, Fx7,
Harvey, HTP, ICS, Jat, MathSAT, Sateen, Simplify, Spear,
STeP, STP, SVC, TSAT, UCLID, Yices, Zap, *Z3 (Microsoft)*

- ▶ SMT-Lib: library of benchmarks

`http://www.smtlib.org`

- ▶ SMT-Comp: annual SMT-Solver competition.

`http://www.smtcomp.org`

Z3: An Efficient SMT Solver

- ▶ *Z3 is a new SMT solver developed at Microsoft Research.*
- ▶ Version 0.1 competed in SMT-COMP'07.
 - ▶ 4 first places.
 - ▶ 7 second places.
- ▶ Version 1.0 was released last week.
- ▶ Free for non-commercial use.
- ▶ Managed (.Net) & Unmanaged (C/C++) APIs are available.
- ▶ `http://research.microsoft.com/projects/z3`
- ▶ More features coming soon.

Applications

- ▶ Test-case generation.
 - ▶ *Pex, SAGE, Yogi, and Vigilante.*
- ▶ Verifying Compiler.
 - ▶ *Spec#/Boogie, VCC/Boogie, and HAVOC/Boogie*
 - ▶ ESC/Java
- ▶ Model Checking & Predicate Abstraction.
 - ▶ *SLAM/SDV and Yogi.*
- ▶ Bounded Model Checking (BMC) & k -induction.
- ▶ Planning & Scheduling.
- ▶ Equivalence checking.

Roadmap

- ▶ Test-case generation
- ▶ Verifying Compiler
- ▶ Model Checking & Predicate Abstraction.
- ▶ Future

Test-case generation

- ▶ Test (correctness + usability) is 95% of the deal:
 - ▶ Dev/Test is 1-1 in products.
 - ▶ Developers are responsible for unit tests.
- ▶ Tools:
 - ▶ Annotations and static analysis (SAL, ESP)
 - ▶ File Fuzzing
 - ▶ *Unit test case generation:*
program analysis tools, automated theorem proving.

Security is Critical

- ▶ Security bugs can be very expensive:
 - ▶ Cost of each MS Security Bulletin: \$600K to \$Millions.
 - ▶ Cost due to worms (Slammer, CodeRed, Blaster, etc.): \$Billions.
 - ▶ *The real victim is the customer.*
- ▶ Most security exploits are initiated via files or packets:
 - ▶ Ex: Internet Explorer parses dozens of files formats.
- ▶ Security testing: *hunting for million-dollar bugs*
 - ▶ Write A/V (always exploitable),
 - ▶ Read A/V (sometimes exploitable),
 - ▶ NULL-pointer dereference,
 - ▶ Division-by-zero (harder to exploit but still DOS attack), ...

Hunting for Security Bugs

- ▶ Two main techniques used by “*black hats*”:
 - ▶ Code inspection (of binaries).
 - ▶ *Black box fuzz testing.*
- ▶ **Black box** fuzz testing:
 - ▶ A form of black box random testing.
 - ▶ Randomly *fuzz* (=modify) a well formed input.
 - ▶ Grammar-based fuzzing: rules to encode how to fuzz.
- ▶ **Heavily** used in security testing
 - ▶ At MS: several internal tools.
 - ▶ Conceptually simple yet effective in practice
 - Has been instrumental in weeding out 1000 of bugs during development and test.*

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Example:

Input x, y

$z = x + y$

If $z > x - y$ Then

 Return z

Else

 Error

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Example:

Input x, y

$z = x + y$

If $z > x - y$ Then

 Return z

Else

 Error

Solve $z = x + y \wedge z > x - y$

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Example:

Input x, y

$z = x + y$

If $z > x - y$ Then

 Return z

Else

 Error

Solve $z = x + y \wedge z > x - y$

$\implies x = 1, y = 1$

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Example:

Input x, y

$z = x + y$

If $z > x - y$ Then

Return z

Else

Error

Solve $z = x + y \wedge \neg(z > x - y)$

Automatic Code-Driven Test Generation

Given program with a set of input parameters.

Generate inputs that maximize code coverage.

Example:

Input x, y

$z = x + y$

If $z > x - y$ Then

Return z

Else

Error

Solve $z = x + y \wedge \neg(z > x - y)$

$\implies x = 1, y = -1$

Method: Dynamic Test Generation

- ▶ *Run* program with *random* inputs.
- ▶ *Collect constraints* on inputs.
- ▶ *Use SMT solver* to generate new inputs.
- ▶ Combination with randomization: DART
(Godefroid-Klarlund-Sen-05)

Method: Dynamic Test Generation

- ▶ *Run* program with *random* inputs.
- ▶ *Collect constraints* on inputs.
- ▶ *Use SMT solver* to generate new inputs.
- ▶ Combination with randomization: DART
(Godefroid-Klarlund-Sen-05)

Repeat while finding new *execution paths*.

DARTish projects at Microsoft

- ▶ *SAGE* (CSE) implements DART for x86 binaries and merges it with “fuzz” testing for finding security bugs.
- ▶ *PEX* (MSR-Redmond FSE Group) implements DART for .NET binaries in conjunction with “parameterized-unit tests” for unit testing of .NET programs.
- ▶ *YOGI* (MSR-India) implements DART to check the feasibility of program paths generated statically using a SLAM-like tool.
- ▶ *Vigilante* (MSR Cambridge) partially implements DART to dynamically generate worm filters.

Initial Experiences with SAGE

25+ security bugs and counting. (most missed by blackbox fuzzers)

- ▶ OS component X

4 new bugs: “This was an area that we heavily fuzz tested in Vista”.

- ▶ OS component Y

Arithmetic/stack overflow in y.dll

- ▶ Media format A

Arithmetic overflow; DOS crash in previously patched component

- ▶ Media format B & C

Hard-to-reproduce uninitialized-variable bug

- ▶ Pex monitors the execution of .NET application using the CLR profiling API.
- ▶ Pex dynamically checks for violations of programming rules, e.g. resource leaks.
- ▶ Pex suggests code snippets to the user, which will prevent the same failure from happening again.
- ▶ *Very instrumental in exposing bugs in .NET libraries.*

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
- ▶ Arithmetic \times Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
 - ▶ Pre-processing step.
 - ▶ Eliminate variables and simplify input formula.
 - ▶ *Significant performance impact.*
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
- ▶ Arithmetic × Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
 - ▶ *Z3 is incremental*: new constraints can be asserted.
 - ▶ **push** and **pop**: (user) backtracking.
 - ▶ Reuse (some) lemmas.
- ▶ “Small models”.
- ▶ Arithmetic × Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
 - ▶ **Given** a set of constraints C , find a model M that *minimizes* the value of the variables x_0, \dots, x_n .
- ▶ Arithmetic \times Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
- ▶ **Given** a set of constraints C , find a model M that *minimizes* the value of the variables x_0, \dots, x_n .
- ▶ **Eager (cheap) Solution:**
Assert C .
While satisfiable
 Peek x_i such that $M[x_i]$ is big
 Assert $x_i < c$, where c is a small constant
Return last found model
- ▶ Arithmetic \times Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
- ▶ **Given** a set of constraints C , find a model M that *minimizes* the value of the variables x_0, \dots, x_n .
- ▶ **Refinement:**
 - ▶ Eager solution stops as soon as the context becomes unsatisfiable.
 - ▶ A “bad” choice (peek x_i) may prevent us from finding a good solution.
 - ▶ Use **push** and **pop** to retract “bad” choices.
- ▶ Arithmetic \times Machine Arithmetic.

Test-case generation & SMT

- ▶ Formulas are usually a big conjunction.
- ▶ Incrementality: solve several similar formulas.
- ▶ “Small models”.
- ▶ Arithmetic \times Machine Arithmetic.
 - ▶ *Precision \times Performance.*
 - ▶ SAGE has flags to abstract expensive operations.

Roadmap

- ▶ Test-case generation
- ▶ Verifying Compiler
- ▶ Model Checking & Predicate Abstraction.
- ▶ Future

The Verifying Compiler

A verifying compiler uses *automated reasoning to check the correctness* of a program that is compiles.

Correctness is specified by *types, assertions, ... and other redundant annotations* that accompany the program.

Hoare 2004

Spec# Approach for a Verifying Compiler

- ▶ *Source Language*

- ▶ C# + goodies = Spec#

- ▶ *Specifications*

- ▶ method contracts,
 - ▶ invariants,
 - ▶ field and type annotations.

- ▶ *Program Logic*

- ▶ Dijkstra's weakest preconditions.

- ▶ *Automatic Verification*

- ▶ type checking,
 - ▶ verification condition generation (VCG),
 - ▶ automatic theorem proving (SMT)

Spec# Approach for a Verifying Compiler

▶ Spec# (annotated C#) \implies Boogie PL \implies Formulas

▶ Example:

```
class C {  
    private int a, z;  
    invariant z > 0  
    public void M()  
        requires a != 0  
        { z = 100/a; }  
}
```

▶ Weakest preconditions:

▶ $wp(S; T, Q) = wp(S, wp(T, Q))$

▶ $wp(\text{assert } C, Q) = C \wedge Q$

Microsoft Hypervisor

- ▶ **Meta OS:** small layer of software between hardware and OS.
- ▶ **Mini:** 60K lines of non-trivial concurrent systems C code.
- ▶ **Critical:** must *guarantee isolation*.
- ▶ **Trusted:** a grand verification challenge.

What is to be verified?

- ▶ Source code: C + x64 assembly.
- ▶ Sample verifiable slices:
 - ▶ **Safety:** Basic memory safety
 - ▶ **Security:** OS isolation
 - ▶ **Utility:** Hypervisor services guest OS with available resources.

Tool: A *Verified C Compiler*

- ▶ VCC translates an *annotated C program* into a *Boogie PL* program.
- ▶ Boogie generates verification conditions:
 - ▶ Z3 (automatic)
 - ▶ Isabelle (interactive)
- ▶ A C-ish memory model
 - ▶ Abstract heaps
 - ▶ Bit-level precision
- ▶ The verification project has very recently started.
- ▶ It is a multi-man multi-year effort.
- ▶ More news coming soon.

Tool: HAVOC

- ▶ A tool for specifying and checking properties of systems software written in C.
- ▶ It also translates annotated C into Boogie PL.
- ▶ It allows the expression of *richer properties about the program heap and data structures* such as linked lists and arrays.
- ▶ HAVOC is being used to specify and check:
 - ▶ Complex locking protocols over heap-allocated data structures in Windows.
 - ▶ Properties of collections such as IRP queues in device drivers.
 - ▶ Correctness properties of custom storage allocators.

Verifying Compilers & SMT

- ▶ *Quantifiers, Quantifiers, ...*
 - ▶ Modeling the runtime.
 - ▶ Frame axioms (“what didn’t change”).
 - ▶ User provided assertions and invariants (e.g., the array is sorted).
 - ▶ Prototyping decision procedures (e.g., reachability, partial orders, ...).
- ▶ Since first-order logic is undecidable, satisfiability is not solvable for arbitrary quantified formulas.
- ▶ Z3: pragmatic approach
 - ▶ *Heuristic Quantifier Instantiation.*
 - ▶ E-matching (i.e., matching modulo equalities).

Heuristic Quantifier Instantiation

- ▶ Semantically, $\forall x_1, \dots, x_n. F$ is equivalent to the infinite conjunction $\bigwedge_{\beta} \beta(F)$.
- ▶ Solvers use heuristics to select from this infinite conjunction those instances that are “relevant”.
- ▶ The key idea is to treat an instance $\beta(F)$ as relevant whenever it contains enough terms that are represented in the solver state.
- ▶ Non ground terms p from F are selected as *patterns*.
- ▶ *E-matching* (matching modulo equalities) is used to find instances of the patterns.
- ▶ Example: $f(a, b)$ matches the pattern $f(g(x), x)$ if $a = g(b)$.

E-matching

- ▶ E-matching is NP-hard.
- ▶ The number of matches can be exponential.
- ▶ It is not refutationally complete.
- ▶ In practice:
 - ▶ Indexing techniques for fast retrieval.
 - ▶ Incremental E-matching.

E-matching: example

- ▶ $\forall x. f(g(x)) = x$
- ▶ Pattern: $f(g(x))$
- ▶ Atoms: $a = g(b), b = c, f(a) \neq c$
- ▶ $\rightarrow \text{instantiate } f(g(b)) = b$

Quantifiers in Z3

- ▶ Z3 uses a E-matching abstract machine.
 - ▶ Patterns \rightsquigarrow code sequence.
 - ▶ Abstract machine executes the code.
- ▶ *Z3 uses new algorithms that identify matches on E-graphs incrementally and efficiently.*
 - ▶ E-matching code trees.
 - ▶ Inverted path index.
- ▶ Z3 garbage collects clauses, together with their atoms and terms, that were useless in closing branches.

Roadmap

- ▶ Test-case generation
- ▶ Verifying Compiler
- ▶ Model Checking & Predicate Abstraction.
- ▶ Future

SLAM: device driver verification

- ▶ <http://research.microsoft.com/slam/>
- ▶ *SLAM/SDV* is a software model checker.
- ▶ Application domain: *device drivers*.
- ▶ Architecture
 - c2bp** C program \rightsquigarrow boolean program (*predicate abstraction*).
 - bebop** Model checker for boolean programs.
 - newton** Model refinement (*check for path feasibility*)
- ▶ SMT solvers are used to perform predicate abstraction and to check path feasibility.
- ▶ c2bp makes several calls to the SMT solver. The formulas are relatively small.

Predicate Abstraction: c2bp

- ▶ **Given** a C program P and $F = \{p_1, \dots, p_n\}$.
- ▶ **Produce** a boolean program $B(P, F)$
 - ▶ Same control flow structure as P .
 - ▶ Boolean variables $\{b_1, \dots, b_n\}$ to match $\{p_1, \dots, p_n\}$.
 - ▶ Properties true of $B(P, F)$ are true of P .
- ▶ Each p_i is a pure boolean expression.
- ▶ Each p_i represents set of states for which p_i is true.
- ▶ Performs modular abstraction.

Abstracting Expressions via F

- ▶ $\text{Implies}_F(e)$

- ▶ Best boolean function over F that implies e

- ▶ $\text{ImpliedBy}_F(e)$

- ▶ Best boolean function over F that implied by e

- ▶ $\text{ImpliedBy}_F(e) = \neg \text{Implies}_F(\neg e)$

Computing $\text{Implies}_F(e)$

- ▶ minterm $m = l_1 \wedge \dots \wedge l_n$, where $l_i = p_i$, or $l_i = \neg p_i$.
- ▶ $\text{Implies}_F(e)$ is the disjunction of all minterms that imply e .
- ▶ Naive approach
 - ▶ Generate all 2^n possible minterms.
 - ▶ For each minterm m , use SMT solver to check validity of $m \implies e$.
- ▶ Many possible optimizations.

Newton

- ▶ Given an error path p in the boolean program B .
- ▶ Is p a feasible path of the corresponding C program?
 - ▶ Yes: found a bug.
 - ▶ No: find predicates that explain the infeasibility.
- ▶ Execute path symbolically.
- ▶ Check conditions for inconsistency using SMT solver.

Roadmap

- ▶ Test-case generation
- ▶ Verifying Compiler
- ▶ Model Checking & Predicate Abstraction.
- ▶ Future

Future work

- ▶ Proof production
- ▶ Interpolants
- ▶ Quantifier elimination
 - ▶ Z3 already supports Fourier-Motzkin elimination.
- ▶ New theories:
 - ▶ Reachability
 - ▶ Sets
 - ▶ Partial orders
 - ▶ Better support for non-linear arithmetic

Proof production

- ▶ It is not required by internal projects.
 - ▶ May be useful in Hypervisor.
- ▶ Requested by potential external users (e.g., Cambridge).
- ▶ *Useful for debugging (and optimizing) Z3.*
- ▶ Our approach:
 - ▶ Goal: reasonable performance when proof production is enabled.
 - ▶ Store a compact representation of deduction steps in a log.
 - ▶ Use a post-processor to transform log into a proof.

Interpolants

- ▶ Requested several times by internal & external users.
- ▶ Useful for predicate abstraction (i.e., predicate discovery).

Conclusion

- ▶ Formal verification is hot at Microsoft.
- ▶ Main applications:
 - ▶ Test-case generation.
 - ▶ Verifying compiler.
 - ▶ Model Checking & Predicate Abstraction.
- ▶ Z3 is a new SMT solver.

<http://research.microsoft.com/projects/z3>